

CS246—Assignment 2 (Spring 2015)

Due Date 1: Wednesday, June 3rd, 04:55pm

Due Date 2: Wednesday, June 10th, 04:55pm

Questions 1a, 2a, 3a, 4a are due on Due Date 1; the remainder of the assignment is due on Due Date 2.

Note: On this and subsequent assignments, you will be required to take responsibility for your own testing. As part of that requirement, this assignment is designed to get you into the habit of thinking about testing *before* you start writing your program. If you look at the deliverables and their due dates, you will notice that there is *no* C++ code due on Due Date 1. Instead, you will be asked to submit test suites for C++ programs that you will later submit by Due Date 2.

Test suites will be in a format compatible with A1Q5 and A1Q6 (as applicable). So if you did a good job writing your `runSuite` script, it will serve you well on this assignment.

Be sure to do a good job on your test suites, as they will be your primary tool for verifying the correctness of your submission. For this reason, **C++ code due on Due Date 2 will only get one release token per day**. We want you to rely on your own pre-written test suite, not Marmoset, to verify correctness.

Note: You must use the C++ I/O streaming and memory management facilities on this assignment. Marmoset will be programmed to **reject** submissions that use C-style I/O or memory management.

Note: Further to the previous note, your solutions may only `#include` the headers `<iostream>`, `<fstream>`, `<sstream>`, `<iomanip>`, and `<string>`. No other standard headers are allowed. Marmoset will check for this.

Note: There will be a handmarking component in this assignment, whose purpose is to ensure that you are following an appropriate standard of documentation and style, and to verify any assignment requirements not directly checked by Marmoset. Please code to a standard that you would expect from someone else if you had to maintain their code. Documentation guidelines have been uploaded to the repository (DocumentationOutline.pdf).

1. In an election, there are n candidates, numbered 1 through n , where n is non-negative. Each voter is allowed one vote for the candidate of their choice. The vote is recorded by number (from 1 to n). The number of voters is unknown beforehand, but is, of course, non-negative. Votes are terminated by end-of-file. Any vote which is not a number from 1 to n is considered an invalid (*spoilt*) vote. The input consists of the names of the candidates (one full name per line), followed by the votes. The first name is considered as candidate 1, and second as candidate 2, and so on. A candidate's name will never contain a numeral.

Write a program to read the data and display the results of the election. For example, given the following data:

Victor Taylor
Denise Duncan

Kamal Ramdhan
Michael Ali
Anisa Sawh
Carol Khan
Gary Owen
3 1 2 5 4 3 5 3 5 3 2 8 1 6 7
7 3 5 6 9 3 4 7 1 2 4 5 5 1 4

your program should produce the following output:

Number of voters: 30
Number of valid votes: 28
Number of spoilt votes: 2

Candidate	Score
Victor Taylor	4
Denise Duncan	3
Kamal Ramdhan	6
Michael Ali	4
Anisa Sawh	6
Carol Khan	2
Gary Owen	3

A starter file `a2q1.cc` has been provided in your `cs246/1155/a2` directory.

- Write the function, `readVotes`, which reads the data from standard input, processes the votes, and records the required information.
- Write the function, `printResults`, which prints the results of the election, using the information obtained by the call to `readVotes` (which you may assume has happened). The data must be presented **exactly** as in the example above. In particular, the Candidate column should be left-aligned and should have a width of 15 characters, and the Score column should be right-aligned with a width of 3 characters (except for the header, which is left-aligned).

Use the following format for the data in the two columns:

- Candidate: left-justified, 15 characters wide
- Score: right-justified, 3 characters wide

Assume there are at most 10 candidates, and each name consists of at least 1 and at most 15 characters (including any spaces).

Note: Do not copy paste the example above to create a test file. The formatting might NOT be correct. This test case has been provided to you (q1.in and q1.out) within the a2 directory

- Due on Due Date 1:** Design a test suite for this program, using the main function provided in the test harness. Call your suite file `suiteq1.txt`. Zip your suite file, together with the associated `.in` and `.out` files, into the file `a2q1.zip`.
- Due on Due Date 2:** Write the program in C++. Save your solution in `a2q1.cc`.

2. In this question we will extend the previous question by implementing **cumulative voting**. In this voting scheme a voter has X number of votes that they can choose to distribute among the candidates which, as before, are numbered from 1 to n . X is an optional positive command line parameter with a default value of n . The input consists of the name of the candidates (one full name per line), followed by some number of lines where each line indicates one voter's distribution of votes, which we call a **ballot**. For a ballot the i^{th} column indicates the number of votes allocated to the i^{th} candidate. A ballot is considered invalid (*spoilt*) if it does not consist of n columns or the sum of the votes in the ballot exceeds X . In addition, the votes allocated to a specific candidate within a ballot are always non-negative.

For example, given the following data:

```
Victor Taylor
Denise Duncan
Kamal Ramdhan
Michael Ali
Anisa Sawh
Carol Khan
Gary Owen
3 0 1 0 0 1 2
1 1 1 1 0 1 2
1 1 1 1 1 1 1
2 1 3 1
7 0 0 0 0 0 0
1 1 1 1 1 1 2
```

your program should produce the following output:

```
Number of voters: 6
Number of valid ballots: 4
Number of spoilt ballots: 2
```

Candidate	Score
Victor Taylor	12
Denise Duncan	2
Kamal Ramdhan	3
Michael Ali	2
Anisa Sawh	1
Carol Khan	3
Gary Owen	5

The formatting requirements and restriction on the number of candidates and their names remains unchanged from the previous question. The program will be run using a test harness equivalent to A1Q6 runSuite.

Note: Look at the sample program `a2/args.cc` to learn how command line arguments are read in by C++ programs.

Note: Do not copy paste the example above to create a test file. The formatting might NOT be correct. This test case has been provided to you (`q2.in` and `q2.out`) within the `a2` directory

- Due on Due Date 1:** Design a test suite (`suiteq2.txt`) for this program. Zip your suite file, together with the associated `.in` and `.out` files, into the file `a2q2.zip`.
- Due on Due Date 2:** Write the program in C++. Save your solution in `a2q2.cc`.

3. Write a program that will process a series of zero or more credit card transactions by reading in the transaction record from standard input, validating the credit card, and then outputting a verdict as to the validity of the credit card. A transaction record consists of the following fields, one per line:

- name on the credit card i.e. first name, optional initial, last name e.g. Jane A. Qiu, Fred Smith
- credit card number e.g. 4556737586899855¹
- card expiry date (1- or 2-digit month [space] 2-digit year), [space], transaction number (positive integer, up to 5 digits in length), [space], date (day [space] month [space] year), time (24-hour time), [space] and amount e.g. 05 18 00001 25 05 2015 2105 45.03

You may assume that all of the information is correctly formatted and is of the appropriate type, so the only form of validation that is required is to check whether or not the credit card number is valid. (Ideally, you'll eventually expand the program to validate the full transaction.) The program will only validate VISA card numbers for now. VISA card numbers start with the digit '4', and are either 13 or 16 digits in length. Note that the number is too large to store into an integer, so you will have to manipulate the string character-by-character, and convert the individual characters to digits as necessary using `stringstream`.

The card number is validated by applying the `Luhn algorithm`. For example, consider the following 16-digit number 4556737586899855.

- (1) If the first digit isn't a 4, the card is invalid.
- (2) If the length of the card number is neither 13 nor 16, the card is invalid.
- (3) The rightmost digit is the **check digit**. In the example, this is the digit 5. Starting from the check digit and moving to the left, double every second digit (marked in bold). The number **4556737586899855** now becomes i.e. 8,5,10,6,14,3,14,5,16,6,16,9,18,8,10,5.
- (4) For each number over 9, add together the digits of the product. The number now becomes 8,5,1+0,6,1+4,3,1+4,5,1+6,6,1+6,9,1+8,8,1+0,5 = 8,5,1,6,5,3,5,5,7,6,7,9,9,8,1,5
- (5) Add all numbers together. The sum is 90.
- (6) Calculate the result of the sum modulo 10 i.e. 90 modulo 10 = 0.
- (7) If the result is 0, the card number is valid.

The format of each record output is as follows:

```
00001 25/05/2015 21:00 $45.03 (4556737586899855, Jane A. Qiu, 05/18) valid[newline]
00001 25/05/2015 21:00 $45.03 (4556737586899856, Fred Smith, 05/18) invalid[newline]
```

Note that the transaction number is specified to have a width of 5 digits, and a fill character of '0' The transaction date has a width of 2 set for the day and the month. The transaction time has a width of 2 for each of the hour and the minutes. The dollar amount has a precision of 2 (hint: investigate use of `fixed` and `showpoint` IO manipulators). The expiry month and year each have a width of 2.

Note: Do not copy paste the example above to create a test file. The formatting might NOT be correct. This test case has been provided to you (q3.in and q3.out) within the a2 directory

As soon as a transaction record's credit card is determined to be invalid, the program stops checking the credit card and moves on to the next record. The program ends upon reaching the end of the input stream.

¹See the file `a2/validCCNums.txt` for a list of 10 valid VISA numbers that you can use in your test cases.

- (a) **Due on Due Date 1:** Design a test suite for this program. Call your suite file `suiteq3.txt`. Zip your suite file, together with the associated `.in` and `.out` files, into the file `a2q3.zip`.
- (b) **Due on Due Date 2:** Write the program in C++. Call your program `validate.cc`.
4. In this question, you will write a program called `ppmTransform` that will read in pixel data for an image from standard input, optionally apply some transformations to the image and write out the final image to standard output in PPM format.

PPM is simple image encoding format. While PPM allows for some variation in format, you must follow the specific formatting rules given in this question. Marmoset will not understand PPM files that do not conform to the given guidelines.

In particular, the output from your program will be a Plain PPM file. Refer to <http://netpbm.sourceforge.net/doc/ppm.html> for additional details on the PPM format (Plain PPM section). We add the following specifications for the image files that should be generated by your program.

The file must start with the following header:

```
P3
Width Height
255
```

P3 is the magic number for a Plain PPM file and 255 is the maximum colour value for the pixel data. These two values should never change for the images your program generates for this assignment. `Width` and `Height` are numbers representing the width and height of the image, respectively. These two values are separated by whitespace.

Each line following the header, should contain the pixel data for a row of the image in order from top to bottom. The pixel data for an entire row should be on a single line, ending with a newline. A single pixel consists of three numbers, between 0 – 255, representing the RGB value for that pixel. Consecutive RGB values should be separated by white space. For a single pixel, the individual RGB values should also be separated by white space. The following is an example image file in the correct format:

```
P3
3 4
255
0 0 0 128 128 128 255 255 255
255 0 0 255 0 0 255 0 0
0 255 0 0 255 0 0 255 0
0 0 255 0 0 255 0 0 255
```

The input to the program will be a sequence of RGB values for an image, the width of the image, the height of the image, followed by a series of transformations to apply to the image. All the input will be separated by whitespace. The possible transformations are ‘‘flip’’, ‘‘rotate’’, and ‘‘sepia’’.

The RGB values for each pixel should be stored in the following structure.

```
struct Pixel {
    unsigned int r; // value for red
    unsigned int g; // value or green
    unsigned int b; // value for blue
};
```

We typically use arrays to store collections of items (say, integers). We can allow for limited growth of a collection by allocating more space than typically needed, and then keeping track of how much space was actually used. We can allow for unlimited growth of the array by allocating the array on the heap and resizing as necessary.

The latter approach will be needed for the image pixel data as the width and height of the image are not known in advance. As the RGB data is read into a one dimensional array, once the width and height are known, the pixel for location (i,j) is accessed via $i * \text{width} + j$.

The following structure encapsulates a partially-filled PPM pixel array:

```
struct PpmArray {
    int size;        // number of pixels the array currently holds
    int capacity;    // number of pixels the array could hold, given current
                    // memory allocation to pixels
    int width;       // width of image
    int height;      // height of image
    Pixel *pixels;
};
```

For memory allocation, you **must** follow this allocation scheme: every `PpmArray` structure begins with a capacity of 0. The first time data is stored in a `PpmArray` structure, it is given a capacity of 5 and space allocated accordingly. If at any point, this capacity proves to be not enough, you must double the capacity (so capacities will go from 5 to 10 to 20 to 40 ...). Note that there is no `realloc` in C++, so doubling the size of an array necessitates allocating a new array and copying items over. Your program must not leak memory.

Your program should support the following image transformations:

- **flip**: Write a function with the following signature

```
void flipHorizontal(PpmArray&)
```

The function modifies the `PpmArray` structure so that it now contains an image that has been flipped horizontally.

- **rotate**: Write a function with the following signature

```
void rotate(PpmArray&)
```

The function modifies the `PpmArray` structure so that it now contains an image that has been rotated 90° clockwise.

If i is the row index and j is the column index of a pixel in the rotated image, then the following will give the required values for the new pixel array, where `ppm` contains the current image.

```
newPixel[i * ppm.height + j] = ppm.pixels[ppm.width * (ppm.height - j - 1) + i];
```

- **sepia**: Write a function with the following signature

```
void sepia(PpmArray&)
```

The function modifies the `PpmArray` structure so that it now contains an image that has a sepia filter applied to the image.

Apply the following formulae (exactly as give) to each pixel in the image. For `Pixel p` with result stored in `Pixel np`:

```
np.r = p.r * .393 + p.g * .769 + p.b * .189
np.g = p.r * .349 + p.g * .686 + p.b * .168
np.b = p.r * .272 + p.g * .534 + p.b * .131
```

If any of the final values above exceed 255, then set them to the maximum value of 255.

You are required to verify that input data conforms to a valid image and valid transformations. For any error conditions, print an informative error message to standard error and abort the program with a nonzero exit status.

The following files are provided for testing (`a2/sample-tests` directory):

- `small.ppm` (original image), `small.in` (program input), `smallFinal.ppm` (transformed image)
- `castle.ppm` (original image), `castle.in` (program input), `castleFinal.ppm` (transformed image). Castle image is from https://commons.wikimedia.org/wiki/File:Castle_Neuschwanstein.jpg and is distributed through the GNU Free Documentation License.

Each sample input file contains the RGB values for the image, followed by the width and height, followed by a list of transformations. Any sample input files you create should also follow this format. Your program should run as follows:

```
./ppmTransform < small.in > output.ppm
```

Make sure the output of your program matches the given output for the sample files. Your mark depends on it.

To view a PPM file on the `linux.student.cs` environment, the following command will work:

```
display file.ppm
```

To view PPM files on your local machine various options exist. Both `ImageMagick` and `gimp` are open-source programs and should work on OS X, Windows or Linux.

A starter file `ppmTransform.cc` has been provided in your `cs246/1155/a2` directory.

- Due on Due Date 1:** Design a test suite for this program. In general, use input/output files of an appropriate size for your test cases. In particular, use `small.in` as the starting point for your test cases where reasonable.
Call your suite file `suiteq4.txt`. Zip your suite file, together with the associated `.in` and `.out` files, into the file `a2q4.zip`.
- Due on Due Date 2:** Write the program in C++. Call your solution `ppmTransform.cc`.